

Topics

Federated Lab

Pytorch

<https://pytorch.org/docs/stable/index.html>

Federated Learning Lab

Based on the work of
Balaji Varatharajan

Reference code and presentation

[GitHub - BalajiAI/Federated-Learning: Implementation of Federated Learning algorithms such as FedAvg, FedAvgM, SCAFFOLD, FedOpt, Mime using PyTorch.](#)

[Aggregation algorithms in Federated Learning](#)

Papers

FedAvg:[\[1602.05629\] Communication-Efficient Learning of Deep Networks from Decentralized Data](#)

FedOpt:(adam,adagrad,yogi):[\[2003.00295\] Adaptive Federated Optimization](#)

SCAFFOLD:[\[1910.06378\] SCAFFOLD: Stochastic Controlled Averaging for Federated Learning](#)

Base Federated Learning Vocabulary

Federated Learning (FL): A distributed learning paradigm where multiple clients (devices) train a shared model without centralizing data.

Client: The individual devices or nodes that perform local training on private data.

Server: The central coordinator that aggregates updates from clients.

Local Update: The process where each client trains the model on its local data.

Global Model: The aggregated model obtained after combining client updates.

Federated Averaging (FedAvg): The baseline aggregation algorithm that averages client model updates.

Communication Round: A complete cycle of local training and subsequent aggregation on the server.

Data Heterogeneity (Non-IID Data): The variability in data distribution across different clients.

Privacy Preservation: Techniques used to ensure client data remains private during training.

Scalability: The system's ability to handle a large number of clients efficiently.

FedAvg

Client i

$$y_i = x$$

$$y_i = y_i - \eta \frac{\partial L}{\partial y_i}$$

communicate y_i

Server

$$x = \frac{1}{|S|} \sum_{i \in S} y_i$$

communicate x

FedAvg

Client i

$$y_i = y_i - \eta \frac{\partial L}{\partial y_i}$$

```
grads = torch.autograd.grad(loss, self.y.parameters())
with torch.no_grad():
    for param, grad in zip(self.y.parameters(), grads):
        param.data = param.data - self.lr * grad.data
```

Server

$$x = \frac{1}{|S|} \sum_{i \in S} y_i$$

```
with torch.no_grad():
    for idx in client_ids:
        for a_y, y in zip(avg_y, self.clients[idx].y.parameters()):
            a_y.data.add_(y.data / int(self.fraction * self.num_clients))
    for param, a_y in zip(self.x.parameters(), avg_y):
        param.data = a_y.data
```

Adaptive Federated Optimization Vocab

Adaptive Learning Rate: An approach where the learning rate is automatically adjusted based on historical gradient information.

FedAdam / FedAdagrad / FedYogi: Variants of adaptive optimizers (inspired by [Adam](#), [Adagrad](#), and [Yogi](#)) tailored for federated settings.

Momentum: A technique that incorporates previous updates to smooth and accelerate convergence.

Bias Correction: Adjustments made (e.g., in Adam) to correct the estimates of moment statistics.

Gradient Scaling: Methods to adjust gradients (or learning rates) based on their magnitudes or variance.

Optimizer Hyperparameters: Parameters such as beta coefficients in Adam that control decay rates and other dynamics.

Convergence Stability: The algorithm's ability to reliably reach a minimum despite data heterogeneity and noisy gradients.

Client Drift: The divergence in local updates due to non-IID data that adaptive methods aim to counteract.

FedAdagrad

Client i

$$y_i = x$$

$$y_i = y_i - \eta_l \frac{\partial L}{\partial y_i}$$

$$\Delta y_i = y_i - x$$

communicate Δy_i

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

$$s = s + g^2$$

$$x = x + \frac{\eta_g}{\sqrt{s + \epsilon}} g$$

communicate x

FedAdagrad

$$\text{Client } i \quad y_i = y_i - \eta_l \frac{\partial L}{\partial y_i}$$

```
with torch.no_grad():
    for param, grad in zip(self.y.parameters(), grads):
        param.data = param.data - self.lr * grad.data
```

$$\Delta y_i = y_i - x$$

```
with torch.no_grad():
    delta_y = [torch.zeros_like(param, device=self.device) for param in self.y.parameters()]
    for del_y, param_y, param_x in zip(delta_y, self.y.parameters(), self.x.parameters()):
        del_y.data += param_y.data.detach() - param_x.data.detach()
    self.delta_y = delta_y
```

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

```
with torch.no_grad():
    for idx in client_ids:
        for grad, diff in zip(gradients, self.clients[idx].delta_y):
            grad.data.add_(diff.data / int(self.fraction * self.num_clients))
```

$$s = s + g^2$$

$$x = x + \frac{\eta_g}{\sqrt{s + \epsilon}} g$$

```
for p,g,s in zip(self.x.parameters(), gradients, self.s):
    s.data += torch.square(g.data)
    p.data += self.lr * g.data / torch.sqrt(s.data + self.epsilon)
```

FedAdam

Client i

$$y_i = x$$

$$y_i = y_i - \eta_l \frac{\partial L}{\partial y_i}$$

$$\Delta y_i = y_i - x$$

communicate Δy_i

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

$$m = \beta_1 m + (1 - \beta_1) g$$

$$v = \beta_2 v + (1 - \beta_2) g^2$$

$$\hat{m} = \frac{m}{1 - \beta_1^t}$$

$$\hat{v} = \frac{v}{1 - \beta_2^t}$$

$$x = x + \eta_g \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$$

communicate x

FedAdam

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

$$m = \beta_1 m + (1 - \beta_1) g$$

$$v = \beta_2 v + (1 - \beta_2) g^2$$

$$\hat{m} = \frac{m}{1 - \beta_1^t}$$

$$\hat{v} = \frac{v}{1 - \beta_2^t}$$

$$x = x + \eta_g \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$$

```
with torch.no_grad():
    for idx in client_ids:
        for grad, diff in zip(gradients, self.clients[idx].delta_y):
            grad.data.add_(diff.data / int(self.fraction * self.num_clients))

for p,g,m,v in zip(self.x.parameters(), gradients, self.m, self.v):
    m.data = self.beta1 * m.data + (1 - self.beta1) * g.data
    v.data = self.beta2 * v.data + (1 - self.beta2) * torch.square(g.data)
    m_bias_corr = #####m#####
    v_bias_corr = #####v#####
    p.data += self.lr * m_bias_corr / (torch.sqrt(v_bias_corr) + self.epsilon)
```

FedYogi

Client i

$$y_i = x$$

$$y_i = y_i - \eta_l \frac{\partial L}{\partial y_i}$$

$$\Delta y_i = y_i - x$$

communicate Δy_i

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

$$m = \beta_1 m + (1 - \beta_1) g$$

$$v = v + (1 - \beta_2) g^2 \odot \text{sgn}(g^2 - v)$$

$$\hat{m} = \frac{m}{1 - \beta_1^t}$$

$$\hat{v} = \frac{v}{1 - \beta_2^t}$$

$$x = x + \eta_g \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$$

communicate x

FedYogi

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

```
with torch.no_grad():
    for idx in client_ids:
        for grad, diff in zip(gradients, self.clients[idx].delta_y):
            grad.data.add_(diff.data / int(self.fraction * self.num_clients))
```

$$m = \beta_1 m + (1 - \beta_1) g$$

$$v = v + (1 - \beta_2) g^2 \odot sgn(g^2 - v)$$

$$\hat{m} = \frac{m}{1 - \beta_1^t}$$

$$\hat{v} = \frac{v}{1 - \beta_2^t}$$

$$x = x + \eta_g \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$$

```
for p,g,m,v in zip(self.x.parameters(), gradients, self.m, self.v):
    m.data = self.beta1 * m.data + (1 - self.beta1) * g.data
    v.data = v.data + (1 - self.beta2) * torch.sign(torch.square(g.data) - v.data) * torch.square(g.data)
    m_bias_corr = #####m_bias_corr#####
    v_bias_corr = #####v_bias_corr#####
    p.data += self.lr * m_bias_corr / (torch.sqrt(v_bias_corr) + self.epsilon)
```

SCAFFOLD Vocab

SCAFFOLD: Stochastic Controlled Averaging for Federated Learning; a method to correct client drift.

Control Variates: Auxiliary variables used to reduce variance in local updates and correct for client drift.

Client Control Variate: A variable maintained at each client to adjust local updates based on estimated drift.

Server Control Variate: The aggregate control variable maintained by the server to guide correction across clients.

Drift Correction: The process of adjusting updates to counter the bias introduced by heterogeneous data distributions.

Variance Reduction: Techniques used to decrease the variability in gradient estimates, enhancing convergence.

Local Gradient Correction: Specific adjustments made to local gradients using control variates.

Stochastic Optimization: The broader framework that underpins methods like SCAFFOLD, dealing with randomness in gradient updates.

Update Correction: The mechanism to adjust the direction and magnitude of client updates based on control variates.

Scaffold(Stochastic Controlled Averaging for Federated Learning)

Client i

$$y_i = x, c = c$$

$$y_i = y_i - \eta_l \left(\frac{\partial L}{\partial y_i} - c_i + c \right)$$

$$c_i^+ = \begin{cases} \text{(i)} \frac{\partial L}{\partial x} & \text{or (ii)} \\ c_i - c_i - c + \frac{1}{K_{\eta_l}}(x - y_i) \end{cases}$$

$$\Delta y_i = y_i - x$$

$$\Delta c_i = c_i^+ - c_i$$

$$c_i = c_i^+$$

communicate $(\Delta y_i, \Delta c_i)$

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

$$\Delta c = \frac{1}{|S|} \sum_{i \in S} \Delta c_i$$

$$x = x + \eta_g g$$

$$c = c + \frac{|S|}{N} \Delta c$$

communicate (x, c)

Scaffold(Stochastic Controlled Averaging for Federated Learning)

Client i

$$y_i = x, c = c$$

$$y_i = y_i - \eta_l \left(\frac{\partial L}{\partial y_i} - c_i + c \right)$$

$$c_i^+ = \begin{cases} \text{(i)} \frac{\partial L}{\partial x} & \text{or (ii)} \\ c_i - c_i - c + \frac{1}{K_{\eta_l}}(x - y_i) \end{cases}$$

$$\Delta y_i = y_i - x$$

$$\Delta c_i = c_i^+ - c_i$$

$$c_i = c_i^+$$

communicate $(\Delta y_i, \Delta c_i)$

```
def client_update(self):
    self.x.to(self.device)
    self.y = deepcopy(self.x) #Initialize local model
    self.y.to(self.device)

    for epoch in range(self.num_epochs):
        data_iter = iter(self.data)
        inputs, labels = next(data_iter)
        inputs, labels = inputs.float().to(self.device), labels.long().to(self.device)
        output = self.y(inputs)
        loss = self.criterion(output, labels)
        grads = torch.autograd.grad(loss, self.y.parameters())

        with torch.no_grad():
            for param, grad, s_c, c_c in zip(self.y.parameters(), grads, self.server_c, self.client_c):
                s_c, c_c = s_c.to(self.device), c_c.to(self.device)
                param.data = param.data - self.lr * (grad.data + (s_c.data - c_c.data))

        if self.device == "cuda": torch.cuda.empty_cache()

    with torch.no_grad():
        delta_y = [torch.zeros_like(param, device=self.device) for param in self.y.parameters()]
        delta_c = deepcopy(delta_y)
        new_client_c = deepcopy(delta_y)

        for del_y, param_y, param_x in zip(delta_y, self.y.parameters(), self.x.parameters()):
            del_y.data += param_y.data.detach() - param_x.data.detach()
        a = (ceil(len(self.data.dataset)) / self.data.batch_size)*self.num_epochs*self.lr
        for n_c, c_l, c_g, diff in zip(new_client_c, self.client_c, self.server_c, delta_y):
            n_c.data += c_l.data - c_g.data - diff.data / a

        for d_c, n_c_l, c_l in zip(delta_c, new_client_c, self.client_c):
            d_c.data.add_(n_c_l.data - c_l.data)

    self.client_c = deepcopy(new_client_c) #Update client_c with new_client_c
    self.delta_y = delta_y
    self.delta_c = delta_c
```

Scaffold(Stochastic Controlled Averaging for Federated Learning)

Server

$$g = \frac{1}{|S|} \sum_{i \in S} \Delta y_i$$

$$\Delta c = \frac{1}{|S|} \sum_{i \in S} \Delta c_i$$

$$x = x + \eta_g g$$

$$c = c + \frac{|S|}{N} \Delta c$$

communicate (x, c)

```
def server_update(self, client_ids):
    self.x.to(self.device)
    for idx in client_ids:
        with torch.no_grad():
            for param, diff in zip(self.x.parameters(), self.clients[idx].delta_y):
                param.data.add_(diff.data * self.lr / int(self.fraction * self.num_clients))
            for c_g, c_d in zip(self.server_c, self.clients[idx].delta_c):
                c_g.data.add_(c_d.data * self.fraction)
```